

DEPARTMENT OF CSE (ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING)

MODULE-01QB: Introduction to Optimization & Mathematical Foundations	
Subject:	NUMERICAL OPTIMIZATION
Sub Code:	23CSE616
Semester:	VI
Department:	CSE (AIML)

Sl. No	Question	CO	Bloom's Level
1.	Formally define the mathematical optimization problem in terms of an objective function and constraints.	CO1	L1 (Remember)
2.	Distinguish between a global minimum and a local minimum in the context of a non-convex objective function.	CO1	L2 (Understand)
3.	Define a linear optimization problem and state the standard form of a linear programming problem.	CO1	L1 (Remember)
4.	Mathematically define a convex set and provide one simple geometric example.	CO1	L1 (Remember)
5.	Write the mathematical definition of the Hessian matrix for a function $\mathbf{f}: \mathbb{R}^n \rightarrow \mathbb{R}$.	CO1	L1 (Remember)
6.	Define the gradient of a scalar-valued function and explain its geometric significance regarding the direction of steepest ascent.	CO1	L2 (Understand)
7.	Calculate the Euclidean norm of the vector $\mathbf{v} = [3, -4]^T$.	CO1	L3 (Apply)
8.	What is the condition for a symmetric matrix to be considered positive definite?	CO1	L1 (Remember)
9.	Define the term "feasible region" in the context of constrained optimization.	CO4	L1 (Remember)
10.	State the multivariable chain rule for a composite function.	CO1	L1 (Remember)
11.	Briefly explain the role of a cost function (or objective function) in an optimization problem.	CO4	L2 (Understand)
12.	How do eigenvalues relate to the definiteness of a Hessian matrix?	CO1	L2 (Understand)

13.	Define unconstrained optimization and state the first-order necessary condition for optimality.	CO4	L1 (Remember)
14.	Define a partial derivative and explain how it differs from a total derivative.	CO1	L2 (Understand)
15.	Briefly define a saddle point in the context of a multivariable function.	CO1	L1 (Remember)

Sl. No	Question	CO	Bloom's Level
-------------------	-----------------	-----------	--------------------------



DEPARTMENT OF CSE (ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING)

1.	Explain the significance of the determinant and trace of a matrix. How do they relate to the eigenvalues?	CO1	L2 (Understand)
2.	Explain the relationship between the gradient of a function and its level sets (contour lines). Illustrate with a diagram.	CO1	L2 (Understand)
3.	Discuss the role of the Hessian matrix in determining the curvature of a function. How does it help in classifying stationary points?	CO1	L2 (Understand)
4.	Compare and contrast Linear Optimization and Nonlinear Optimization. Provide a mathematical formulation for each.	CO4	L2 (Understand)
5.	Describe the $\ \cdot\ _1$, $\ \cdot\ _2$, $\ \cdot\ _\infty$ and norms. Explain their geometric interpretations (unit circles) and where each might be used.	CO1	L2 (Understand)
6.	Derive or explain the First-Order Necessary Condition (FONC) and Second-Order Sufficient Condition (SOSC) for local minimizers in unconstrained optimization.	CO1	L2 (Understand)
7.	Elaborate on the differences between convex and non convex optimization problems. Why is convexity desirable in optimization?	CO4	L4 (Analyze)
8.	Define the Jacobian matrix for a vector-valued function. Explain how it differs from the Hessian matrix.	CO1	L2 (Understand)
9.	Explain the Taylor series expansion for a multivariable function up to the second order. How is this used in optimization?	CO1	L2 (Understand)
10.	Explain the difference between equality and	CO4	L2

	inequality constraints. How do they alter the search space?		(Understand)
11.	What does it mean for two vectors to be orthogonal? Explain the concept of an orthonormal basis.	CO1	L2 (Understand)
12.	Prove that the gradient points in the direction of the steepest ascent of the function.	CO1	L3 (Apply)
13.	Discuss why Linear Algebra is considered the backbone of modern optimization and Machine Learning. Give specific examples of matrix operations used in ML.	CO5	L2 (Understand)
14.	Classify stationary points into minima, maxima, and saddle points using eigenvalues of the Hessian.	CO1	L3 (Apply)
15.	Explain the concept of Eigen decomposition and its relevance in analyzing the geometry of quadratic functions.	CO1	L2 (Understand)

Sl. No	Question	CO	Bloom's Level
1.	A manufacturing company wants to minimize transportation costs while satisfying demand at three	CO2	L3 (Apply)



DEPARTMENT OF CSE (ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING)

	different retail locations from two warehouses. Formulate this as a Linear Programming problem, defining decision variables, objective function, and constraints.		
2.	You are a financial analyst trying to maximize returns on a portfolio of stocks while minimizing risk (variance). Describe how you would formulate this as a constrained optimization problem. Identify the vector/matrix components involved.	CO5	L4 (Analyze)
3.	A nutritionist needs to design a diet plan that meets minimum daily requirements for Vitamin A, B, and C using three food items, each with different costs and nutrient contents. Formulate the mathematical optimization model to minimize the cost.	CO2	L3 (Apply)
4.	A cloud computing provider must allocate CPU and RAM to incoming tasks to maximize total throughput without exceeding physical server limits. Formulate this scenario as an integer or linear optimization problem.	CO2	L3 (Apply)

5.	Explain how finding the shortest path in a navigation app (like Google Maps) can be viewed as an optimization problem. Define the nodes, edges, costs, and the objective.	CO4	L4 (Analyze)
6.	A factory produces two products, P1 and P2. P1 requires 2 hours of labor and 3 units of raw material. P2 requires 4 hours of labor and 1 unit of raw material. Total labor available is 100 hours, and raw material is 80 units. Formulate the profit maximization problem.	CO2	L3 (Apply)
7.	Given a set of data points ($(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$), formulate the Ordinary Least Squares (OLS) regression problem as an unconstrained optimization problem using vector calculus notation.	CO5	L3 (Apply)
8.	An e-commerce warehouse needs to decide how much stock to order to minimize storage and ordering costs while preventing stockouts. Formulate the objective function involving holding costs and shortage penalties.	CO4	L3 (Apply)
9.	A marketing team has a fixed budget to spend across Social Media, TV, and Print ads. Each channel has a different "Return on Ad Spend" (ROAS) and saturation point. Formulate an optimization problem to maximize total revenue.	CO2	L3 (Apply)
10.	Engineers must design a bridge truss to minimize weight while supporting a specific load without yielding. Describe the decision variables (e.g., cross-sectional areas) and constraints (stress limits).	CO4	L3 (Apply)
11.	An electrical grid needs to balance power generation from wind, solar, and coal to match demand at minimum cost. Formulate this as a constrained optimization problem, considering the variability of renewables.	CO2	L4 (Analyze)
12.	Describe the geometric intuition of finding the "best" separating hyperplane between two classes of data points. How is this an optimization problem?	CO5	L2 (Understand)
13.	A robot arm needs to move from point A to B minimizing energy. Constraints include joint limits and obstacle avoidance. Formulate the path planning optimization problem.	CO2	L3 (Apply)

14.	An airline needs to assign crews to flights to minimize layover costs and ensure legal rest periods. Describe the complexity of this problem and why it is a constrained optimization task.	C04	L4 (Analyze)
15.	A data center wants to minimize cooling energy consumption while maintaining safe server temperatures. Formulate the relationship between fan speed (variable) and temperature (constraint).	C04	L3 (Apply)



DEPARTMENT OF CSE (ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING)

Sl. No	Question	CO	Bloom's Level
1.	Imagine a neural network's loss function as a physical terrain. Explain the concepts of "valleys" (minima), "peaks" (maxima), and "flat regions" (plateaus). How does the gradient guide us through this terrain?	C05	L4 (Analyze)
2.	Why is it mathematically easier to train a Linear Regression model (convex loss) compared to a Deep Neural Network (non-convex loss)? Discuss the implications for finding the global minimum.	C05	L5 (Evaluate)
3.	Intuitively explain how adding an ℓ_1 norm (Lasso) vs. an ℓ_2 norm (Ridge) to the loss function affects the model weights. Why does ℓ_1 lead to sparsity?	C05	L4 (Analyze)
4.	In Deep Learning, what is the "exploding gradient" problem? Explain this using the concept of the Chain Rule and repeated matrix multiplication.	C05	L2 (Understand)
5.	How does the curvature of the loss function (represented by the Hessian) affect the speed of learning? Why do we struggle in areas of high curvature?	C05	L4 (Analyze)
6.	Explain how the eigenvalues of the Hessian matrix at a critical point determine the stability of that point in a machine learning model's training process.	C01	L5 (Evaluate)
7.	Discuss the Bias-Variance tradeoff from an optimization perspective. Does minimizing training error (optimization) always lead to lower test error (generalization)?	C05	L5 (Evaluate)
8.	In high-dimensional ML problems, saddle points are more common than local minima. Explain why this happens and why standard gradient descent might get stuck or slow down near them.	C05	L4 (Analyze)

9.	Why does input feature scaling (normalization) speed up the convergence of gradient descent? Explain using the geometry of the cost function (elliptical vs. spherical contours).	C05	L4 (Analyze)
10.	Optimization algorithms are often treated as black boxes in ML libraries (like scikit-learn). Explain the dangers of this approach without understanding the underlying calculus.	C04	L5 (Evaluate)
11.	What does it mean for a matrix to be ill-conditioned in the context of Linear Regression? How does this affect the stability of the model solution?	C01	L4 (Analyze)
12.	Why can't we initialize all weights in a neural network to zero? Explain using the concept of symmetry breaking and gradients.	C05	L2 (Understand)
13.	Discuss whether getting stuck in a local minimum is always a bad thing in Deep Learning. Are all local minima created equal?	C05	L5 (Evaluate)
14.	Explain how the "Curse of Dimensionality" affects optimization. As we add more features (dimensions), what happens to the volume of the search space?	C05	L2 (Understand)
15.	Explain overfitting as "over-optimizing" on the training set. How do we mathematically constrain the optimization process to prevent this (e.g., early stopping)?	C05	L4 (Analyze)